

# Lesson 6 Ornithology

by Lorraine Gaudio

Lesson generated on September 21, 2025



**BOISE STATE UNIVERSITY**

# Contents

1	📌 Welcome back to R!	2
2	Quick Warm-up	3
3	Build a Data Frame	5
4	Bracket Syntax	7
5	Subsetting	8
6	Extraction	11
6.1	\$ vs. [] . . . . .	11
6.2	Selecting Columns . . . . .	12
7	Insertion	14
7.1	Overwrite Values . . . . .	14
8	Renaming Columns	15
9	Replace Missing Values	16
10	📌 Assignment	17
10.1	Task 0 . . . . .	17
10.2	Task 1 . . . . .	17
10.3	Task 2 . . . . .	18
10.4	Task 3 . . . . .	18
10.5	Task 4 . . . . .	18
10.6	Task 5 . . . . .	19
10.7	Task 6 . . . . .	19
11	Save and Upload	20
12	Today you practiced:	21

# 1. ☒ Welcome back to R!

In lesson one, through four you learned about objects, vectors, functions and how to handle missing data in R. In this lesson, you will learn about subsetting, extraction, and insertion in data frames. You rarely analyse an entire data table at once. You nearly always pull out specific rows or columns, or write back cleaned-up values. Mastering R's bracket "[ ]" syntax is therefore foundational.

To begin Lesson 6, follow these steps:

1. Open your course project for RStudio
2. Create a new file. Today, let's try ☐ "R Markdown" (File > New File > R Markdown).
3. Type in the code provided in this document as you follow along with the video. Pause the video at anytime to answer assignment questions, dig deeper or add memo notes.

## Lesson Overview

By the end of Lesson 6 you will be able to:

1. ☐ Remember – State R's basic subsetting template: `x[rows, cols]`.
2. ☐ Understand – Describe the difference between `$` and `[ ]` for data frames.
3. ☐ Apply – Use logical tests + `which()` to select rows.
4. ☐ Analyze – Summarize a subset to answer a question.
5. ☐ Evaluate – Choose an insertion (overwrite) vs. safe copy strategy

Keep these goals in mind as you move through each section.

## 2. Quick Warm-up

Type the following code in a new code chunk and run.

```
# Review the `sample()` function from lesson 3.
set.seed(1)

args(sample) # □ Check the arguments of sample()
```

```
## function (x, size, replace = FALSE, prob = NULL)
## NULL
```

Type the following code in a new code chunk and run.

```
trial_vec <- sample(x = c(1:3, NA, NA, NA, NA), size = 25, replace = TRUE)
trial_vec
```

```
## [1] 1 NA NA 1 2 NA NA 3 NA 2 3 3 1 NA NA 2 NA NA 2 NA 1 NA NA NA 1
```

Type the following code in a new code chunk and run. Create a memo note, demonstrate learning skill(s) used.

```
trial_vec <- sample(size = 25, replace = TRUE, x = c(1:3, NA, NA, NA, NA))
trial_vec
```

```
## [1] 1 NA NA NA 2 2 NA 1 NA NA 1 NA 3 NA 2 2 NA NA NA NA NA 2 NA 1 NA
```

□ NOTICE: How does changing the order of arguments affect the output? Create a memo note, demonstrate learning skill(s) used.

Review `is.na()` and `which()` from lesson 4.

```
# Review is.na() and which() from lesson 4 and Nesting from lesson 3
is.na(trial_vec)
```

```
## [1] FALSE TRUE TRUE TRUE FALSE FALSE TRUE FALSE TRUE TRUE FALSE TRUE
## [13] FALSE TRUE FALSE FALSE TRUE TRUE TRUE TRUE TRUE FALSE TRUE FALSE
## [25] TRUE
```

```
which(is.na(trial_vec))
```

```
## [1] 2 3 4 7 9 10 12 14 17 18 19 20 21 23 25
```

```
which(!is.na(trial_vec))
```

```
## [1] 1 5 6 8 11 13 15 16 22 24
```

□ Are you ready? Remember, you can summon help whenever you need it.

### 3. Build a Data Frame

In lesson 2, you learned how to create a data frame using the function `data.frame()` to combine three vectors of equal length. You created these vectors using the `c()` function. In this lesson, you will create the vectors using the `ifelse()` function. Type the following code in a new code chunk.

```
?ifelse
```

```
## starting httpd help server ... done
```

```
args("ifelse") # Check the arguments of ifelse()
```

```
## function (test, yes, no)  
## NULL
```

`ifelse()` is a vectorized function that allows you to create a new variable based on a condition. Type the following code in a new code chunk.

```
# Example: Create a new vector trial_vec2 based on trial_vec  
trial_vec2 <- ifelse(trial_vec == 2, "two", "not two")
```

This creates a new vector where each element is “two” if the corresponding element in `trial_vec` is 2, and “not two” otherwise. It takes three arguments:

- condition: a logical test (e.g., `trial_vec == 2`)
- value\_if\_true: the value to assign if the condition is TRUE
- value\_if\_false: the value to assign if the condition is FALSE

Now, let’s create a data frame using `trial_vec` and `trial_vec2`. Type the following code in a new code chunk and run.

```
# Create a data frame with trial_vec and trial_vec2  
trial_df <- data.frame(trial_vec, trial_vec2)
```

```
# View(trial_df) # Check the data frame with two columns  
View(trial_df) # (opens a new tab in RStudio)
```

□ That was a practice round. Now, let's create a more complex data frame with two long vectors of equal length.

#### Field Biology: Bird Banding Observations □

This theme simulates data from a field biologist observing a population of birds, where most are untagged, but a small number have been previously captured and marked with colored leg bands. This is a very common practice in ecology and behavioral studies.

We'll simulate a small inventory called `field_sightings` with two columns:

- `Status` – Describes whether the sighted bird is “Untagged” (the common case) or “Tagged” (the rare case).
- `Band_Color` – If a bird is untagged, the value is “None.” If it is tagged, the value is one of several possible band colors.

Type the following code in a new code chunk.

```
set.seed(8)
Status <- sample(c(rep("Untagged", 3091), rep("Tagged", 287)))
Band_Color <- ifelse(Status == "Untagged", "None",
  sample(c("Red", "Blue", "Green", "Yellow",
    "Silver", "Black", "Orange",
    "White"), length(Status), replace = TRUE))
field_sightings <- data.frame(Status, Band_Color)

# View the data
View(field_sightings)
```

□ Look deeper: Why does the code have line breaks with indentation? What is the advantage? Explain when to break and how far to indent. Create a memo note, demonstrate learning skill(s) used.

## 4. Bracket Syntax

□ Review Lesson 2 about indexing vectors `x[]`. Create a memo note, demonstrate learning skill(s) used. Create a memo note, demonstrate learning skill(s) used.

In this lesson we will use the same syntax to subset data frames.

- `x` is the object (vector, data frame, etc.)
- `[]` is the bracket operator for subsetting
- `x[rows, cols]` extracts specific rows and columns from `x`.
- `rows` and `cols` can be: indices (1, 3:5), names ("Band\_Color"), or a logical vector / expression (`Status == "Tagged"`).

Type the following code in a new code chunk and run.

```
field_sightings[1:3, ] # □ Shows full row extraction
```

Status	Band_Color
Untagged	None
Untagged	None
Tagged	White

*# R's subsetting format: `x[rows, columns]`*

```
field_sightings[1:3, 1] # □ Shows first three rows of first column
```

```
## [1] "Untagged" "Untagged" "Tagged"
```

```
field_sightings[1:3, c(1, 2)] # □ Shows first three rows of first and second
```

Status	Band_Color
Untagged	None
Untagged	None
Tagged	White



## 5. Subsetting

We'll start with `trial_df` from earlier. Type the following code in a new code chunk and run.

```
trial_df[trial_df$trial_vec == 2, ]
```

	trial_vec	trial_vec2
NA	NA	NA
NA.1	NA	NA
NA.2	NA	NA
5	2	two
6	2	two
NA.3	NA	NA
NA.4	NA	NA
NA.5	NA	NA
NA.6	NA	NA
NA.7	NA	NA
15	2	two
16	2	two
NA.8	NA	NA
NA.9	NA	NA
NA.10	NA	NA
NA.11	NA	NA
NA.12	NA	NA
22	2	two
NA.13	NA	NA
NA.14	NA	NA

```
# Test direct subsetting (will include NA rows)  
direct_subset <- trial_df[trial_df$trial_vec == 2, ]  
print(direct_subset)
```

```
##      trial_vec trial_vec2  
## NA         NA    <NA>  
## NA.1        NA    <NA>  
## NA.2        NA    <NA>  
## 5           2     two  
## 6           2     two  
## NA.3        NA    <NA>  
## NA.4        NA    <NA>
```

```
## NA.5      NA      <NA>
## NA.6      NA      <NA>
## NA.7      NA      <NA>
## 15        2      two
## 16        2      two
## NA.8      NA      <NA>
## NA.9      NA      <NA>
## NA.10     NA      <NA>
## NA.11     NA      <NA>
## NA.12     NA      <NA>
## 22        2      two
## NA.13     NA      <NA>
## NA.14     NA      <NA>
```

```
trial_df[trial_df$trial_vec == 2, ]
```

	trial_vec	trial_vec2
NA	NA	NA
NA.1	NA	NA
NA.2	NA	NA
5	2	two
6	2	two
NA.3	NA	NA
NA.4	NA	NA
NA.5	NA	NA
NA.6	NA	NA
NA.7	NA	NA
15	2	two
16	2	two
NA.8	NA	NA
NA.9	NA	NA
NA.10	NA	NA
NA.11	NA	NA
NA.12	NA	NA
22	2	two
NA.13	NA	NA
NA.14	NA	NA

□ Identify: which rows are extracted?

- TRUE (2 == 2)
- FALSE (1 == 2)
- FALSE (3 == 2)
- NA (NA == 2)

□ **Goal: isolate only the tagged birds.**

□ **Step 1** – create a logical test: `[]` vs. `$`. Type the following code in a new code chunk and run.

```
field_sightings[field_sightings$Status == "Tagged", ]
```

`field_sightings[field_sightings$Status == "Tagged", ]` returns the rows where the condition is TRUE, along with all columns.

```
field_sightings$Status == "Tagged"      # TRUE/FALSE for each row
```

`field_sightings$Status == "Tagged"` returns a logical vector.

□ **Step 2** – Store subset of `field_sightings` data frame as a new data frame. Type the following code in a new code chunk and run.

```
tagged_only <- field_sightings[field_sightings$Status == "Tagged", ]
```

□ Look deeper: How many *rows* did we keep? How could you check the number of rows in a data set? Create a memo note, demonstrate learning skill(s) used, including method you used to answer the question.

Type the following code in a new code chunk and run.

```
n_manual <- nrow(tagged_only) # □ Check number of rows  
nrow(tagged_only)
```

```
## [1] 287
```

## 6. Extraction

Using `which()` is a strategy to handle potential NA values in your data. When the test could include NA, wrap it in `which()` to drop unknowns.

Type the following R script in your document script and run.

```
rows_without_na <- which(trial_df$status == 2)

which_subset <- trial_df[rows_without_na, ]

print(which_subset)
```

```
## trial_vec trial_vec2
## 5      2      two
## 6      2      two
## 15     2      two
## 16     2      two
## 22     2      two
```

□ Explore and Play: What is the difference between running `print(which_subset)` and `which_subset`? Create a memo note, demonstrate learning skill(s) used.

When you use a direct logical condition like `trial_df$status == 2`, any NA values in the Status column will result in NA in the logical vector. You get only the indices where the condition is TRUE, effectively dropping/ignoring any rows where the comparison resulted in NA. You get a result that excludes rows with NA in the Status column

### 6.1 \$ vs. [ ]

Type the following code in a new code chunk and run.

```
row_idx <- which(field_sightings["Status"] == "Tagged" & field_sightings["Band_Color"] == "Blue")
```

The bracket notation `["column_name"]` uses matrix-style indexing.

You can access the same columns using `$` notation where `field_sightings$status` replaces `field_sightings["Status"]` and `field_sightings$Band_Color` replaces `field_sightings["Band_Color"]`. Type the following code in a new code chunk and run.

```
row_idx2 <- which(field_sightings$Status == "Tagged" & field_sightings$Band_Color == "Blue")
```

□ Common Pain Point: Using `which()` but forgetting to subset columns. Type the following code in a new code chunk and run.

```
subset_blue <- field_sightings[row_idx, ]  
head(subset_blue) # □ Practice Task 5 logic with known data
```

	Status	Band_Color
182	Tagged	Blue
208	Tagged	Blue
293	Tagged	Blue
319	Tagged	Blue
502	Tagged	Blue
635	Tagged	Blue

```
head(x = subset_blue, n = 3)
```

	Status	Band_Color
182	Tagged	Blue
208	Tagged	Blue
293	Tagged	Blue

□ Explain the function `head()` and how it works with the `n` argument. Create a memo note, demonstrate learning skill(s) used.

## 6.2 Selecting Columns

□ Method 1: Select columns by name using indexing.

Recall that columns can be chosen by index or name.

- By index: `field_sightings[, 2]` (second column)
- By name : `field_sightings[, "Band_Color"]`

```
head(subset_blue[, "Band_Color"])
```

```
## [1] "Blue" "Blue" "Blue" "Blue" "Blue" "Blue"
```

You can combine row + column ideas: band color of the first 10 birds. Type the following code in a new code chunk and run.

```
field_sightings[1:10, "Band_Color"]
```

```
## [1] "None" "None" "White" "None" "None" "None" "None" "None" "None"
## [10] "None"
```

□ Why did we use `head()` in the first chunk and not in the second? Create a memo note, demonstrate learning skill(s) used.

□ Method 2: Select columns by name using `$`. Type the following code in a new code chunk and run.

```
Untagged_Birds <- field_sightings$Band_Color[field_sightings$Status == "Untagged"]
```

This creates a vector called `Untagged_Birds` that contains only the band colors from rows where the status is "Untagged". Type the following code in a new code chunk and run.

```
Tagged_Birds <- field_sightings$Band_Color[field_sightings$Status == "Tagged"]
```

It's particularly useful when you want to:

- Analyze the distribution of "Band\_Colors" within each "Status".
- Compare properties between different categories.
- Create separate vectors for further analysis by category.

This creates a vector called `Tagged_Birds` with only band colors from tagged birds rows.

□ Bonus: Two conditions + column subset

Use `field_sightings`: Pick band colors that are White and Status = Tagged. Type the following code in a new code chunk and run.

```
tagged_data <- field_sightings[field_sightings$Status == "Tagged", ]
```

```
white_rows <- which(tagged_data$Band_Color == "White")
```

```
white_tagged <- tagged_data[white_rows, ]
```

## 7. Insertion

Insertion is the process of adding or modifying values in a data frame. You can insert new values or overwrite existing ones.

### 7.1 Overwrite Values

Suppose we decide to rename “Untagged” to “Wild\_Hatch” in the Status column.

□ Good practice: make a backup first! Type the following code in a new code chunk and run.

```
reintroduction <- field_sightings      # copy
indices_untagged <- which(field_sightings$Status == "Untagged")
reintroduction[indices_untagged, "Band_Color"] <- "Wild_Hatch"
```

□ Check-in: Look at the format of that code: `x[row, column] <- value`. The `<-` operator is not just for assigning a value to a variable name. `<- Wild_Hatch` is assigned as the new value "Wild\_Hatch" to replace the existing value "Untagged". The column name remains "Status".

```
# Verify the change
unique(reintroduction$Band_Color)
```

```
## [1] "Wild_Hatch" "White"    "Orange"    "Silver"    "Red"
## [6] "Green"      "Yellow"    "Blue"      "Black"
```

□ Comment: Why is `field_sightings` unchanged while editing `reintroduction`? Why is this good practice? Create a memo note, demonstrate learning skill(s) used.

## 8. Renaming Columns

Example 1: To rename a column, you can use the `colnames()` function. Type the following code in a new code chunk and run.

```
names(reintroduction)

## [1] "Status" "Band_Color"

colnames(reintroduction) <- c("Band_Status", "Bird_Age")
head(reintroduction)
```

Band_Status	Bird_Age
Untagged	Wild_Hatch
Untagged	Wild_Hatch
Tagged	White
Untagged	Wild_Hatch
Untagged	Wild_Hatch
Untagged	Wild_Hatch

□ Identify: Which columns are renamed? □ Look deeper: Create a list of all the different ways you might identify the column names of a data frame? Create a memo note, demonstrate learning skill(s) used.

Example 2: You can also rename a specific column by index. Type the following code in a new code chunk and run.

```
reintroduction2 <- field_sightings # □ Always copy first
names(reintroduction2)[which(names(reintroduction2) == "Status")] <- "Band_Status"
names(reintroduction2)[which(names(reintroduction2) == "Band_Color")] <- "Bird_Age"
```



## 9. Replace Missing Values

In lesson 4, you learned about handling missing values. Now, let's practice replacing missing values in a data frame. Recall that when a numeric vector has gaps, you can impute with the mean.

Type the following code in a new code chunk and run.

```
# Before
trial_vec

## [1] 1 NA NA NA 2 2 NA 1 NA NA 1 NA 3 NA 2 2 NA NA NA NA NA 2 NA 1 NA

trial_vec[is.na(trial_vec)] <- mean(trial_vec, na.rm = TRUE)
```

The <- operator is not just for assigning a value to a variable name. Here the <- operator is used to assign the mean of the non-NA values‘

**Left side of the assignment:** trial\_vec[is.na(trial\_vec)]

- x[] index of a vector: Recall we can use “:” like 1:3
- index value is is.na(trial\_vec) making logical vector (TRUE/FALSE) marking which elements are NA.

**Right side of the assignment:** mean(trial\_vec, na.rm = TRUE)

- mean() calculates the mean of the non-NA values in trial\_vec

The assignment operator <- takes the value from the right side, places it into the specific positions identified on the left side.

```
# After
trial_vec

## [1] 1.0 1.7 1.7 1.7 2.0 2.0 1.7 1.0 1.7 1.7 1.0 1.7 3.0 1.7 2.0 2.0 1.7 1.7 1.7
## [20] 1.7 1.7 2.0 1.7 1.0 1.7
```

R took all the non-missing values, calculated their mean and replaced all the NA values with this decimal mean

## 10. ☒ Assignment

Now it's your turn to practice creating and using vector objects. Follow the tasks below to complete part of the **technical skill practice assignment**.

1. Work through each task in order. Replace the \_\_\_\_ placeholder with your code or short written answer.
2. Run each completed line to be sure no errors appear and objects show in the Environment.
3. When finished, save your workspace and submit this R Markdown file (RMD) plus the .RData file.

### 10.1 Task 0

☐ Setup: Load the built-in data set mtcars and take a quick look. Identify what the columns contain and meaning of that content.

```
data("mtcars") # already in memory but this keeps the workflow explicit
?mtcars       # help file for variable descriptions
View(mtcars)  # spreadsheet view (optional)
```

☐ Comment: Explain R's basic subsetting template. Hint: It looks like data[ rows , columns ]

### 10.2 Task 1

☐ Manual vs. Automatic

1. Create Manual\_Cars: all rows where am == 1 (manual), all columns.

```
Manual_Cars <- ____
```

2. Use nrow() to record how many manual cars there are (store in n\_manual).

```
n_manual <- ____ # numeric count
```

## 10.3 Task 2

### □ MPG by Cylinder

Extract mpg for each cylinder group into separate objects. Hint: Repeat the bracket pattern

```
Four_Cyl_MPG <- mtcars[mtcars$cyl == 4, "mpg"]
Six_Cyl_MPG <- ____
Eight_Cyl_MPG <- ____
```

## 10.4 Task 3

### □ Analyze Fuel Efficiency

1. Compute the mean mpg for each object.

```
mean_4 <- mean(____)
mean_6 <- mean(____)
mean_8 <- mean(____)
```

2. □ Comment: Which cylinder group is most fuel-efficient? Answer in one word or number.

Most efficient: \_\_\_\_

## 10.5 Task 4

### □ Safe vs. Direct Editing

1. Make a safe copy of mtcars called cars\_safe.

```
cars_safe <- ____
```

2. Rename the 'am' column in cars\_safe to "Transmission\_Type".

```
colnames(cars_safe)[____] <- ____
```

3. □ Comment: Explain *why* working on cars\_safe (a copy) is safer than editing mtcars directly.

Answer: \_\_\_\_

## 10.6 Task 5

### □ Practice which()

Using `cars_safe`, extract the rows for V-shaped engines (`vs == 0`) *and* horsepower (`hp`) above 150. Use `which()` for the row indices and keep only `mpg`, `hp`, and `cyl` columns. Name the object `hi_power_v`.

1. Create `row_idx` for the condition above.

- `cars_safe$vs == 0`
- `&`
- `cars_safe$hp > 150`
- `cars_safe[row_idx, c("mpg", "hp", "cyl")]`

```
row_idx <- ____
```

2. Subset `cars_safe` with `row_idx` and columns `mpg`, `hp`, `cyl`. Name object `hi_power_v`.

```
hi_power_v <- ____  
head(hi_power_v)
```

## 10.7 Task 6

### □ Insertion Example

Add a new column called "Efficiency\_Class" to `cars_safe`. Label "High" if `mpg >= 25`, "Low" otherwise.

Hint: use `ifelse()` and `cars_safe$mpg >= 25, "?", "?"`

```
cars_safe$Efficiency_Class <- ____
```

## 11. Save and Upload

1. You will be submitting **both** the R Markdown and the workspace file. The workspace file saves all the objects in your environment that you created in this lesson. You can save the workspace by running the following command in a code chunk of the R Markdown document:

```
save.image("Assignment6_Workspace.RData")
```

Or you can click the “Save Workspace” button in the Environment pane.

□ **Always save the R documents before closing.**

2. Find the assignment in this week’s module in Canvas and upload **both** the RMD and the workspace file.

## 12. Today you practiced:

- Reading R's subsetting template `x[rows, cols]`.
  - Extracting rows with logical tests and `which()`.
  - Selecting columns by index or name.
  - Safely overwriting values (insertion) after making a copy.
  - Renaming columns and imputing simple missing values.
- Great job! Subsetting is the gateway to every data-cleaning task.